

SVO2-SAM3 Analyzer

Software Solution Plan

Multi-Camera SVO2 Processing with SAM 3 Object Detection

Stereo RGB + Depth + IMU | KITTI-Style Annotations

Version: **1.0**

Date: **January 2026**

Author: **Mauricio Mesones**

Classification: **Internal**

Executive Summary

This document outlines the software architecture for a web-based application that processes Stereolabs ZED 2i camera recordings (SVO2 files) using Meta's Segment Anything Model 3 (SAM 3) for comprehensive object detection and analysis. The solution extracts and correlates stereo RGB images (left/right), depth maps, and IMU sensor data, maintaining full traceability from raw SVO2 frames through to KITTI-style annotation outputs.

100% LOCAL DEPLOYMENT

All models, inference, and data processing run entirely on the local workstation.
No cloud services, external APIs, or internet connection required for operation.

Target Hardware:

- Intel Core Ultra 9 processor
- 128GB DDR5 RAM
- NVIDIA GeForce RTX 5090 (32GB VRAM)
- Ubuntu 22.04.5 LTS (64-bit, GNOME 42.9)

SVO2 File Contents

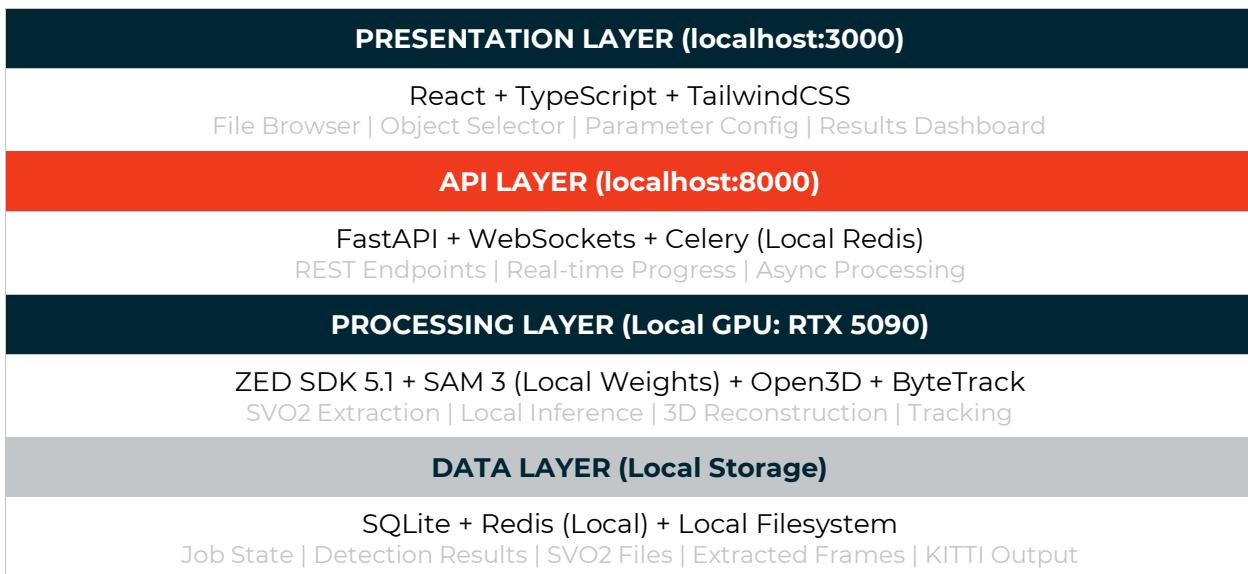
Each SVO2 file contains synchronized multi-modal sensor data that must be extracted and correlated:

| Data Stream | Description |
|----------------------------|---|
| Left RGB | Primary camera image (used for SAM 3 inference) |
| Right RGB | Secondary stereo camera image (for disparity/3D reconstruction) |
| Depth Map | 32-bit float depth per pixel in millimeters (Neural mode) |
| Point Cloud | XYZRGBA per pixel for 3D spatial data |
| IMU - Accelerometer | Linear acceleration (m/s^2) in X, Y, Z axes |
| IMU - Gyroscope | Angular velocity (rad/s) in X, Y, Z axes |
| IMU - Orientation | Quaternion (w, x, y, z) from sensor fusion |
| Timestamps | Nanosecond-precision timestamps for frame/IMU sync |

System Overview

The SVO-SAM3 Analyzer is a full-stack web application consisting of three primary layers: a React-based frontend for user interaction, a FastAPI backend for orchestration and processing, and a processing engine that handles SVO extraction and SAM 3 inference.

High-Level Architecture (Local)



Web UI Components

1. Directory Explorer & File Selector

The file browser provides a familiar tree-view interface for navigating the server's file system and selecting SVO files for processing.

- Tree-view directory navigation with expandable folders and breadcrumb path display
- Multi-select capability with checkboxes for batch processing multiple SVO files
- File filtering (.svo2 extension) with search functionality
- File metadata preview (size, duration, resolution, FPS, IMU rate, recording date)
- Drag-and-drop support for reordering selected files
- Configurable root directories (restricted to authorized paths)

2. Object Class Selector

SAM 3 supports text-prompt segmentation for any describable concept. The selector provides both preset categories and custom prompt entry.

| Category | Available Objects |
|-----------------------|--|
| Vehicles | Haul truck, Excavator, Loader, Light vehicle, Water truck, Dozer |
| Personnel | Person, Worker with PPE, Worker without PPE |
| Equipment | GET (Ground Engaging Tool), Bucket, Tire, Cable, Hose |
| Infrastructure | Berm, Road edge, Stockpile, Building, Power line |
| Safety Items | Cone, Barrier, Sign, Light tower |
| Custom | User-defined text prompts (free-form entry) |

- Toggle switches for each object class with select all/none options
- Custom prompt text field for detecting any describable object
- Save/load preset configurations for common detection scenarios
- Confidence threshold slider per class (0.0 - 1.0)

3. SAM 3 Configuration Panel

| Parameter | Default | Description |
|---------------------|------------|---|
| Detection Threshold | 0.5 | Minimum confidence score for valid detections |
| Mask Threshold | 0.5 | Binary mask threshold for segmentation output |
| Frame Sampling | 1 | Process every Nth frame (1 = all frames) |
| Max Detections | 100 | Maximum objects per frame |
| NMS IoU Threshold | 0.7 | Non-maximum suppression overlap threshold |

| | | |
|-----------------|------------------|---|
| Model Variant | sam3-base | Model size: tiny, small, base, large |
| Precision Mode | FP16 | Inference precision: FP32, FP16, BF16 |
| Batch Size | 8 | Frames processed per GPU batch |
| Enable Tracking | Yes | Enable ByteTrack for object persistence |
| Export 3D Data | Yes | Include 3D bounding boxes in output |

4. Processing Controls

- Start Processing button (enabled when files selected and at least one object class enabled)
- Real-time progress bar with estimated time remaining
- Per-file progress indicators for batch processing
- Live detection count and GPU utilization display
- Pause/Resume and Cancel controls
- Processing log viewer with error highlighting

5. Results Dashboard

Upon completion, the dashboard presents comprehensive analysis results with interactive visualizations.

- Summary statistics: total objects, unique tracks, processing time, frames analyzed
- Object count breakdown by class (bar chart visualization)
- Timeline view showing object presence over video duration
- Interactive data table with sorting, filtering, and search
- Frame viewer with overlay visualization of detections
- 3D point cloud viewer for spatial analysis (using Three.js)
- Export options: KITTI dataset (.zip), JSON, CSV, COCO format, annotated video

Processing Pipeline

Stage 1: SVO2 Extraction & Frame Registry

The ZED SDK extracts all sensor streams from SVO2 files and registers each frame with a unique identifier for full traceability through the annotation pipeline.

- Open SVO2 file with Neural depth mode for AI-enhanced depth estimation
- Extract stereo RGB pair (LEFT and RIGHT views) as PNG or numpy arrays
- Extract depth map (32-bit float, millimeters) aligned to left camera
- Extract point cloud (XYZRGBA) for 3D reconstruction
- Query IMU data: accelerometer (m/s^2), gyroscope (rad/s), orientation quaternion
- Record nanosecond timestamps for frame-IMU synchronization
- Generate unique frame_id (format: {svo2_hash}_{frame_number}) for correlation

Frame Registry Schema

Each extracted frame is registered with complete metadata for annotation traceability:

| Field | Type | Description |
|----------|--------|--|
| frame_id | string | Unique identifier: {svo2_hash}_{frame_num} |

| | | |
|--------------------------|-----------------|---|
| svo2_file | string | Source SVO2 filename |
| svo2_frame_idx | integer | Frame index within SVO2 file |
| timestamp_ns | integer | Nanosecond timestamp from camera |
| timestamp_iso | string | ISO 8601 formatted timestamp |
| left_rgb_path | string | Path to left camera RGB image |
| right_rgb_path | string | Path to right camera RGB image |
| depth_path | string | Path to depth map file (.npy or .png) |
| pointcloud_path | string | Path to point cloud file (.ply) |
| imu_accel | float[3] | Accelerometer [x, y, z] in m/s ² |
| imu_gyro | float[3] | Gyroscope [x, y, z] in rad/s |
| imu_orientation | float[4] | Quaternion [w, x, y, z] |
| camera_intrinsics | object | {fx, fy, cx, cy, k1, k2, p1, p2} |
| camera_extrinsics | object | Stereo baseline and rotation |

Stage 2: SAM 3 Local Inference

SAM 3 runs entirely on the local RTX 5090 GPU. Model weights are loaded from disk once at startup and cached in VRAM for the duration of processing.

- Load SAM 3 weights from `~/.cache/jebi/models/sam3/` (no network calls)
- Apply `torch.compile` optimization for RTX 5090 Blackwell architecture
- Batch frames according to configured batch size (8-16 for 32GB VRAM)
- Execute local GPU inference with text prompts for all enabled object classes
- Post-process on GPU: apply confidence threshold, NMS, generate instance masks
- Output: per-frame list of detections with masks, boxes, scores, and class labels

Stage 3: 3D Reconstruction

- Project 2D segmentation masks onto depth data using camera intrinsics
- Generate per-object point clouds using Open3D
- Compute oriented bounding boxes (center, dimensions, rotation)
- Calculate physical measurements: volume, surface area, distance from camera

Stage 4: Object Tracking

- ByteTrack associates detections across frames using IoU and appearance features
- Assign persistent track IDs to maintain object identity
- Handle occlusion and re-identification with configurable track buffer
- Generate trajectory data for each tracked object

Technology Stack (Local Deployment)

All components run locally on the workstation. No external services or cloud dependencies are required after initial setup.

| Layer | Technology | Purpose |
|------------|------------------------------|--|
| Frontend | React 18 + TypeScript | Component framework (local dev server) |
| Frontend | TailwindCSS | Styling with Jebi brand colors |
| Frontend | React Query | Server state management |
| Frontend | Three.js | 3D point cloud visualization |
| Backend | FastAPI | REST API (localhost:8000) |
| Backend | WebSockets | Real-time progress updates |
| Backend | Celery + Redis | Local async task queue |
| Backend | SQLAlchemy | ORM for local database |
| AI/ML | PyTorch 2.7+ (local) | Deep learning framework |
| AI/ML | SAM 3 weights (local) | Downloaded model checkpoint (~2.4GB) |
| AI/ML | CUDA 12.6 + cuDNN | Local GPU acceleration |
| Processing | ZED SDK 5.1 (pyzed) | SVO2 file processing |
| Processing | Open3D | 3D geometry processing |
| Processing | ByteTrack | Multi-object tracking |
| Data | SQLite / PostgreSQL | Local persistent storage |
| Data | Redis (local) | Caching and task broker |
| Data | Local filesystem | SVO2 files and outputs |

Local Model & Environment Setup

One-time setup downloads all required model weights and dependencies. After initial setup, the system operates fully offline.

SAM 3 Model Weights

| Model Variant | Size | VRAM Required | Recommended For |
|-------------------|---------|---------------|----------------------|
| sam3-tiny | ~400 MB | 4 GB | Quick testing |
| sam3-small | ~900 MB | 8 GB | Balanced performance |
| sam3-base | ~1.8 GB | 12 GB | Production (default) |
| sam3-large | ~2.4 GB | 16 GB | Maximum accuracy |

Local Storage Paths

- Model weights: `~/.cache/jebi/models/sam3/` (downloaded once, ~2.4GB)
- ZED SDK: `/usr/local/zed/` (system installation)
- Processing cache: `~/.cache/jebi/svo2_cache/` (extracted frames)
- Output directory: `~/jebi_outputs/` (KITTI datasets)
- Database: `~/.local/share/jebi/db.sqlite` (job history, settings)
- Redis data: `/var/lib/redis/` (task queue state)

GPU Configuration

RTX 5090 optimization settings for maximum local inference throughput:

| Setting | Configuration |
|-----------------------|---|
| CUDA Version | 12.6+ (local installation) |
| cuDNN Version | 9.x (local installation) |
| PyTorch Build | <code>torch==2.7+</code> with CUDA 12.6 support |
| Precision Mode | <code>BF16 (bfloat16)</code> for RTX 5090 Blackwell |
| torch.compile | <code>mode='reduce-overhead'</code> for inference |
| Memory Format | <code>channels_last</code> for CNN optimization |
| Batch Size | 8-16 frames (adjustable based on VRAM) |
| TF32 Math | Enabled for matrix operations |

Offline Operation Guarantee

- No HuggingFace Hub calls after initial model download
- No telemetry or analytics transmitted
- All inference runs on local GPU (RTX 5090)
- Web UI served from localhost (127.0.0.1:3000)
- API backend on localhost (127.0.0.1:8000)

- Can operate in air-gapped environments after setup

API Endpoints

File Management

| Method | Endpoint | Description |
|-------------|----------------------------|-----------------------------|
| GET | /api/files/browse | List directory contents |
| GET | /api/files/metadata/{path} | Get SVO file metadata |
| POST | /api/files/validate | Validate SVO file integrity |

Processing Jobs

| Method | Endpoint | Description |
|-------------|------------------------|-----------------------------|
| POST | /api/jobs/create | Create new processing job |
| GET | /api/jobs/{id} | Get job status and progress |
| POST | /api/jobs/{id}/start | Start processing |
| POST | /api/jobs/{id}/pause | Pause processing |
| POST | /api/jobs/{id}/cancel | Cancel processing |
| GET | /api/jobs/{id}/results | Get processing results |
| WS | /ws/jobs/{id}/progress | Real-time progress stream |

Configuration

| Method | Endpoint | Description |
|-------------|----------------------------|-------------------------------|
| GET | /api/config/object-classes | List available object classes |
| GET | /api/config/presets | List saved presets |
| POST | /api/config/presets | Save new preset |
| GET | /api/config/model-info | Get SAM 3 model details |

Export

| Method | Endpoint | Description |
|-------------|------------------------|------------------------------------|
| GET | /api/export/{id}/kitti | Export KITTI-format dataset (.zip) |
| GET | /api/export/{id}/json | Export full results as JSON |
| GET | /api/export/{id}/csv | Export summary as CSV |
| GET | /api/export/{id}/coco | Export COCO-format annotations |
| POST | /api/export/{id}/video | Generate annotated video overlay |

Output: KITTI-Style Annotations

The primary output follows the KITTI 3D Object Detection benchmark format, extended with SVO2 correlation metadata. This enables direct use with autonomous driving toolchains while maintaining full traceability to source data.

KITTI Label File Format

Each frame generates a .txt label file with one line per detected object:

| Column | Index | Description |
|-------------------|-------|--|
| type | 0 | Object class (e.g., HaulTruck, Person, Excavator) |
| truncated | 1 | Truncation level: 0.0 (fully visible) to 1.0 (truncated) |
| occluded | 2 | Occlusion state: 0=visible, 1=partial, 2=mostly, 3=unknown |
| alpha | 3 | Observation angle relative to camera center $[-\pi, \pi]$ |
| bbox | 4-7 | 2D bounding box: left, top, right, bottom (pixels) |
| dimensions | 8-10 | 3D dimensions: height, width, length (meters) |
| location | 11-13 | 3D center location: x, y, z in camera coords (meters) |
| rotation_y | 14 | Rotation around Y-axis in camera coordinates $[-\pi, \pi]$ |
| score | 15 | Detection confidence score (0.0 to 1.0) |

Extended Metadata File

Each frame also generates a json metadata file linking annotations to SVO2 source data:

| Field | Type | Description |
|--------------------------------|----------------|--|
| frame_id | string | Unique frame identifier for SVO2 correlation |
| svo2_source | object | {file, frame_idx, timestamp_ns} |
| stereo_images | object | {left_path, right_path} |
| depth_source | object | {path, format, unit: 'mm'} |
| imu_data | object | {accel, gyro, orientation, temp} |
| camera_calibration | object | Full intrinsics and extrinsics |
| detections | array | Extended detection data per object |
| detections[].track_id | integer | Persistent tracking ID across frames |
| detections[].mask_rle | string | Run-length encoded segmentation mask |
| detections[].pointcloud | object | {count, centroid, obb_rotation} |

Output Directory Structure

The KITTI-style output follows a standardized directory layout:

- image_2/ — Left camera RGB images (PNG, numbered 000000.png, 000001.png, ...)
- image_3/ — Right camera RGB images (PNG, same numbering)
- depth/ — Depth maps (16-bit PNG or .npy float32)
- velodyne/ — Point clouds in .bin format (KITTI compatible)
- label_2/ — KITTI format label files (.txt)
- metadata/ — Extended JSON metadata with SVO2 correlation
- calib/ — Camera calibration files (P0, P1, P2, P3, R0_rect, Tr_velo_cam)
- oxts/ — IMU/GPS data in KITTI oxts format
- frame_registry.json — Master index mapping frame numbers to SVO2 source

Implementation Phases

| Phase | Deliverables |
|-------|---|
| 1 | Project setup, FastAPI skeleton, React scaffold, database schema |
| 2 | File browser UI, directory API, SVO2 metadata extraction |
| 3 | SAM 3 local model integration, inference pipeline, GPU optimization |
| 4 | Object selector UI, configuration API, preset management |
| 5 | 3D reconstruction, ByteTrack integration, tracking pipeline |
| 6 | Results dashboard, visualizations, KITTI export functionality |
| 7 | WebSocket progress, error handling, performance tuning |
| 8 | Testing, documentation, local deployment packaging |

Next Steps

1. Review and approve this software solution plan
2. Set up RTX 5090 workstation with Ubuntu 22.04.5 LTS (64-bit, GNOME 42.9)
3. Install NVIDIA drivers, CUDA 12.6, and cuDNN 9.x locally
4. Install ZED SDK 5.1 (local license) and verify camera connectivity
5. Download SAM 3 model weights (one-time, ~2.4GB) for local storage
6. Define initial set of object classes for mining operations
7. Provide sample SVO2 files for development and testing
8. Verify air-gapped operation capability (disconnect network, run inference)